

Reinforcement Drive

The Race For AI





Table of contents



Introduction

Tools and Technologies

Reinforcement Learning – PPO

Why not SAC?

Tuning

Results

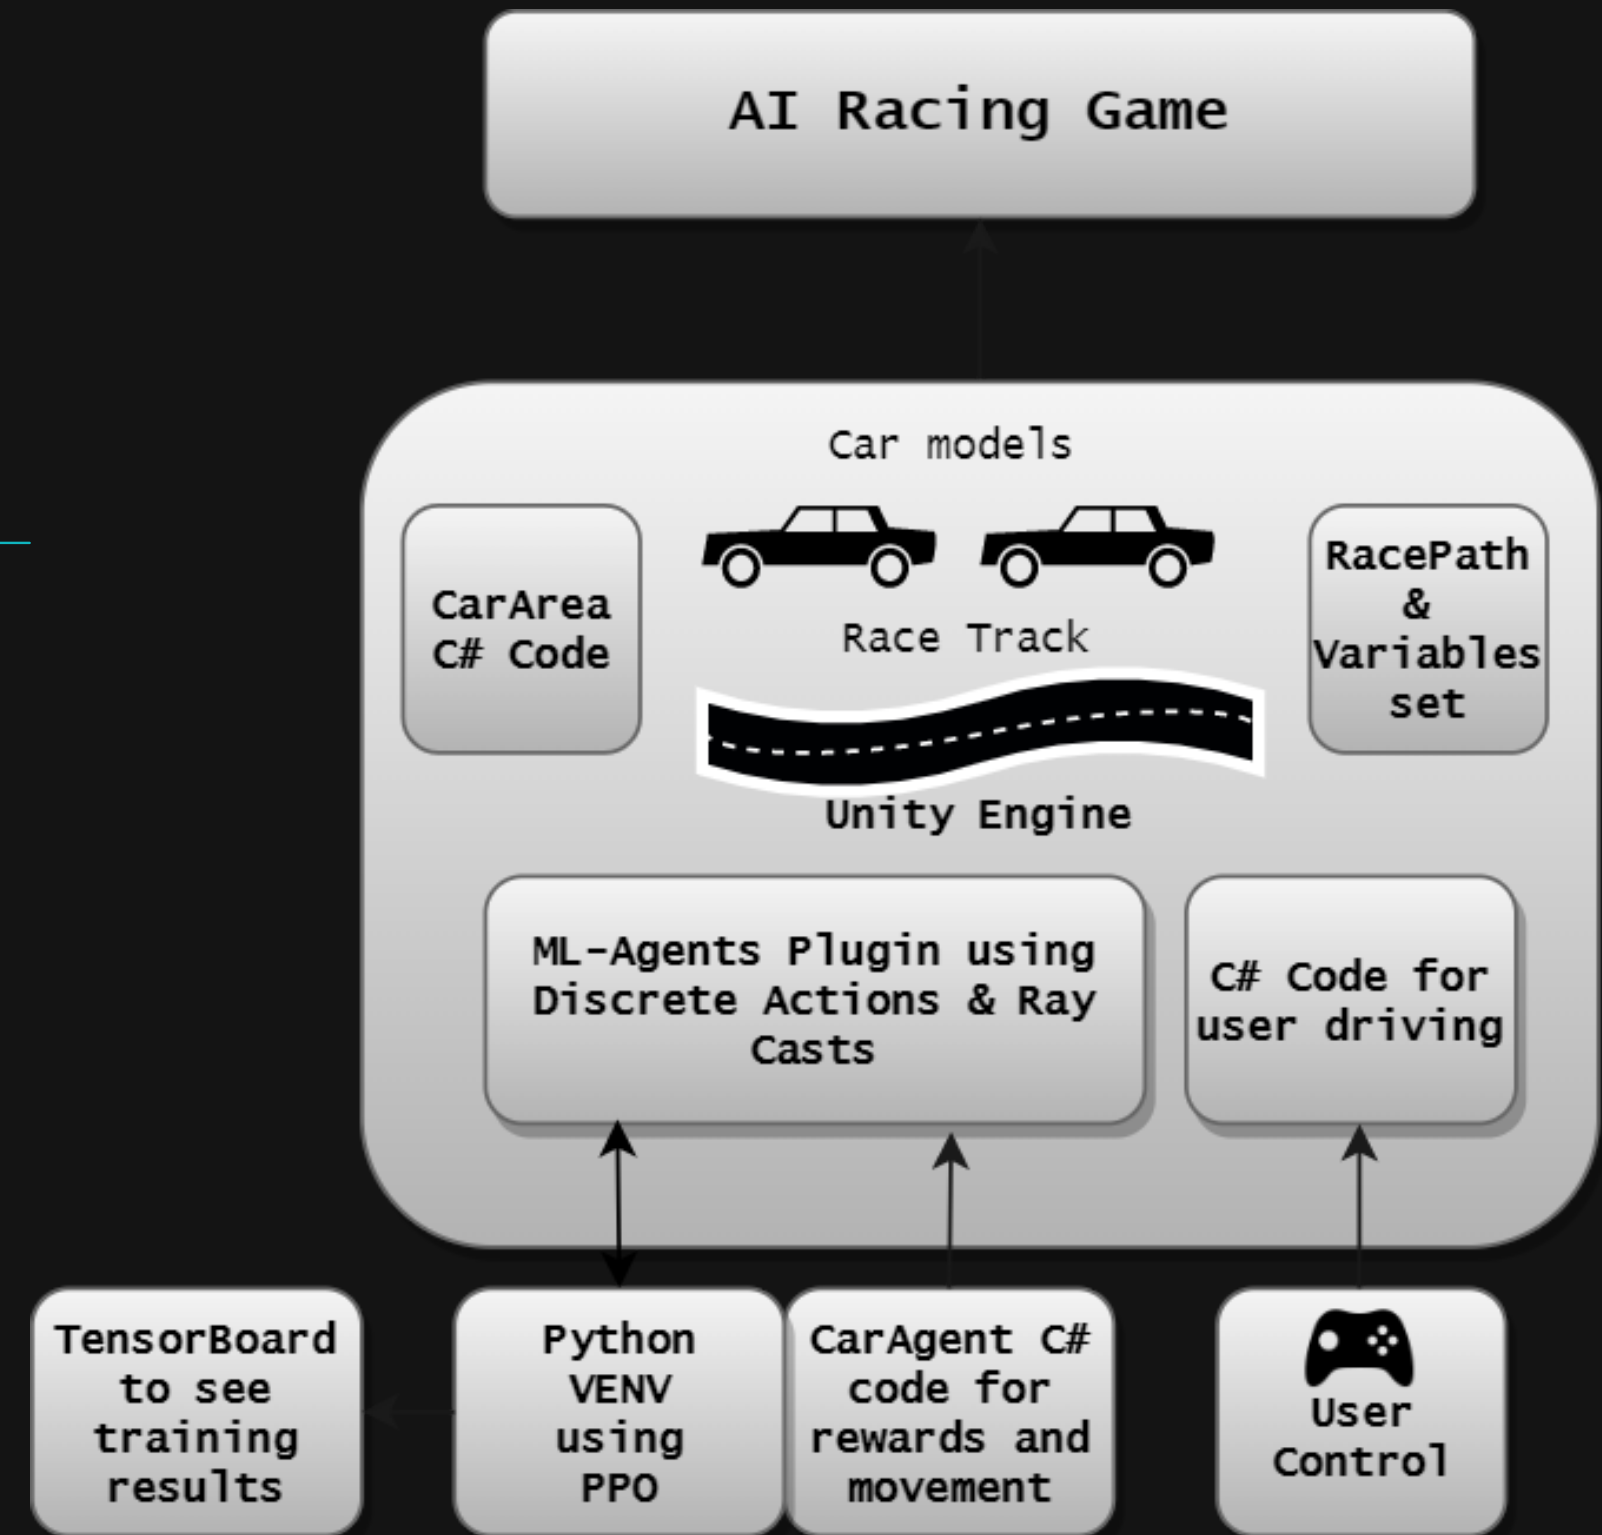
Conclusion

Introduction

- Goal – Create a racing game with AI-controlled opponents using reinforcement learning
- Develop an understanding of AI and reinforcement learning concepts
- Implement key game mechanics – AI behaviour, track creation, car physics

Technologies and Tools

- Unity game engine
- C# programming language
- ML-Agents plugin for Unity
- Proximal Policy Optimization (PPO)
- Python Virtual Environment



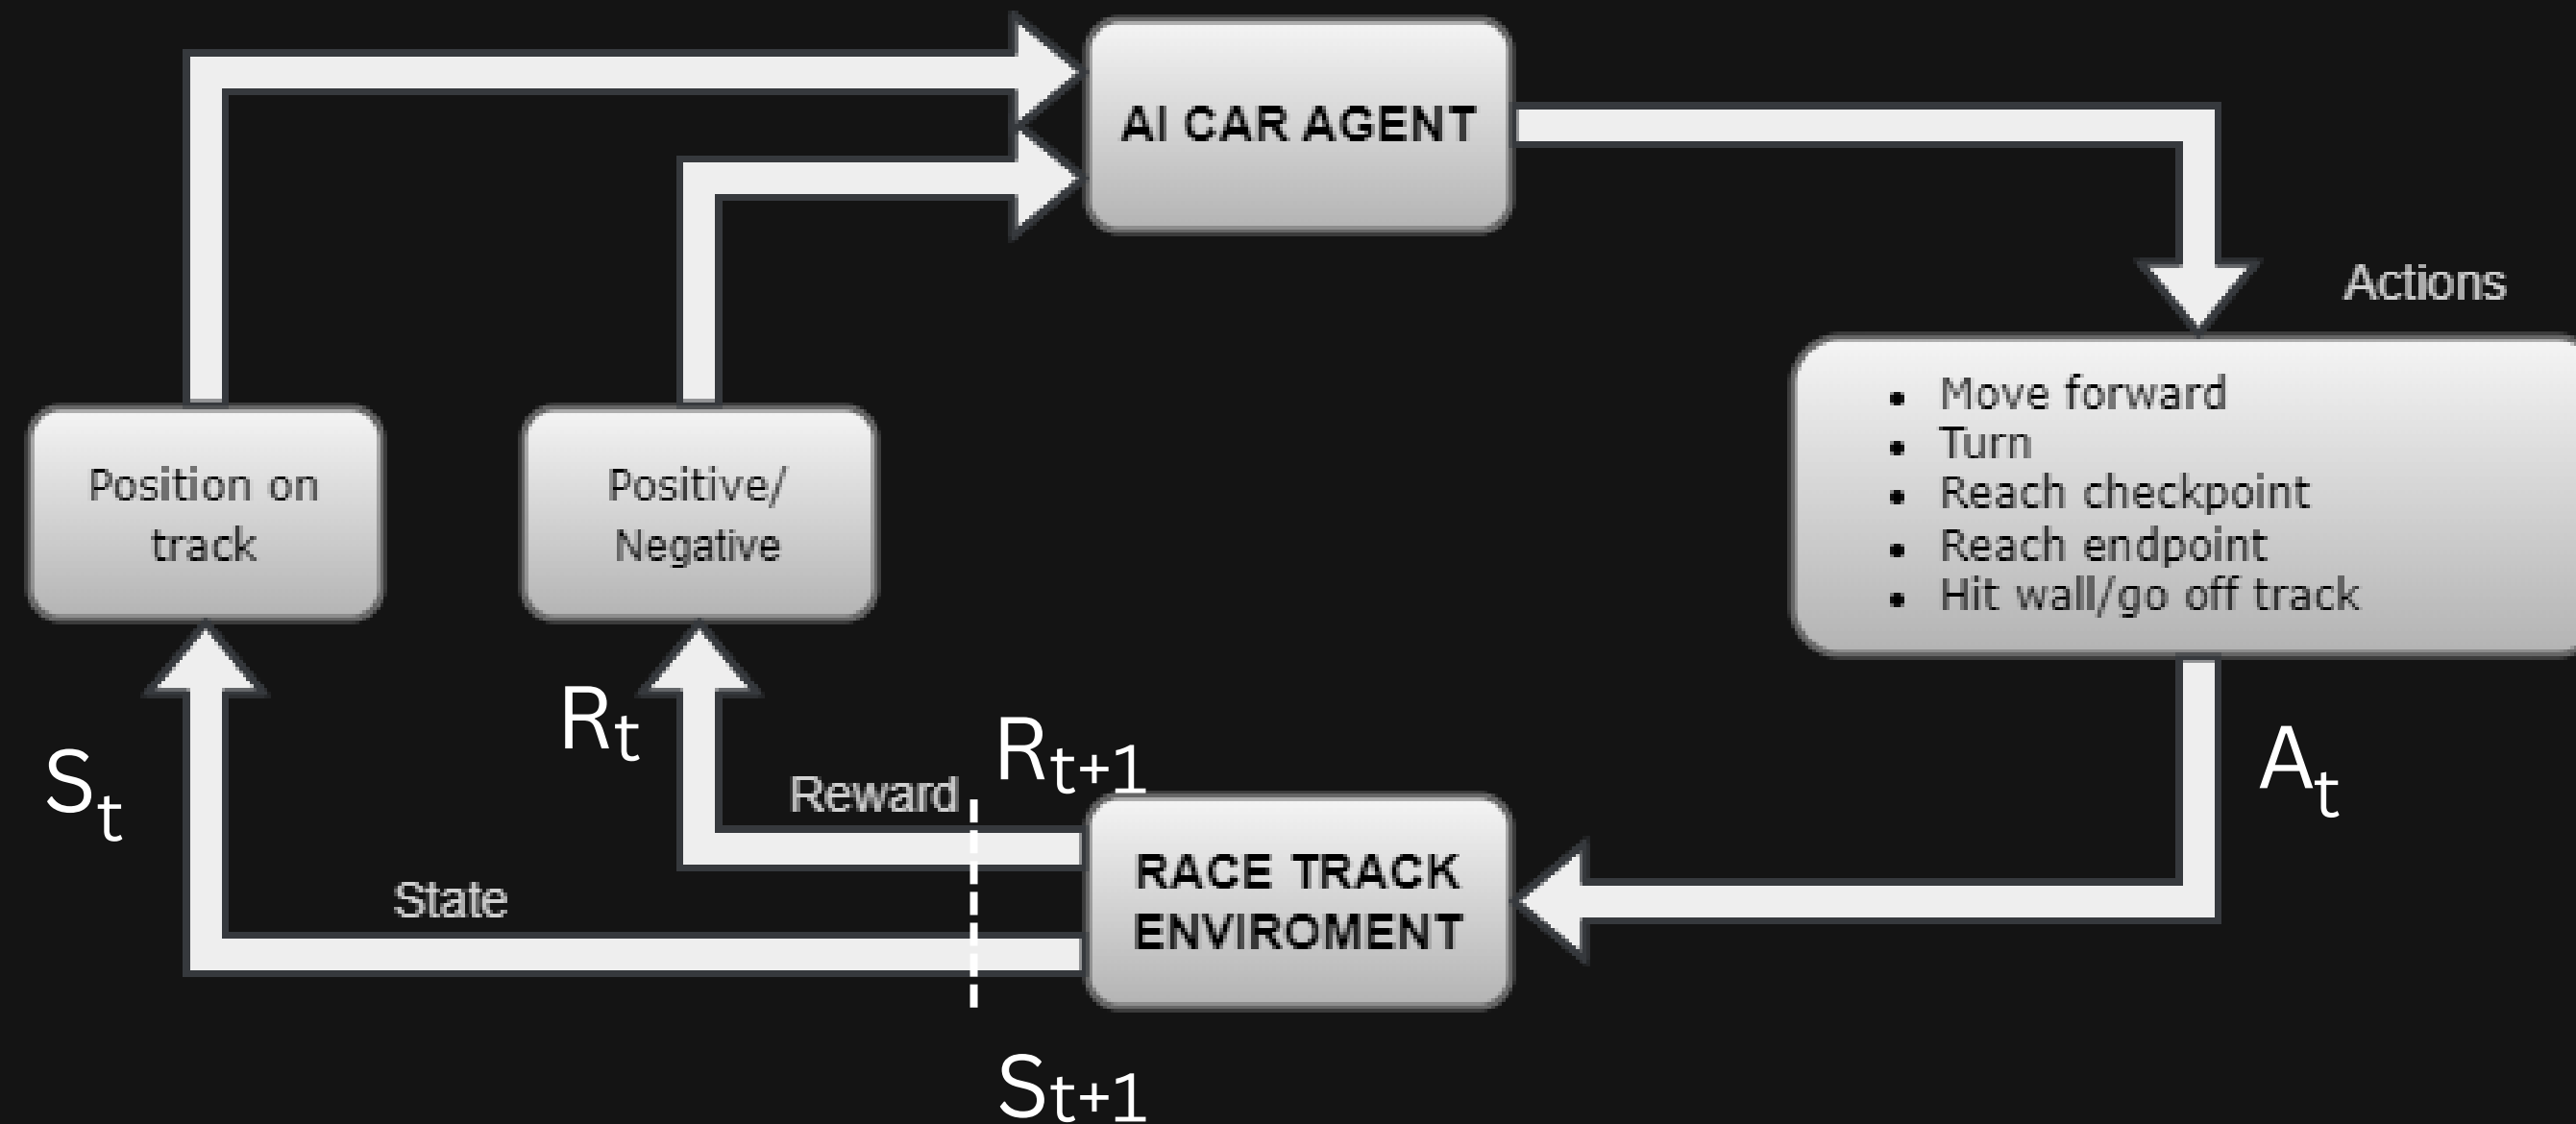
Reinforcement Learning – PPO

- Reinforcement learning: The agent learns by interacting with the environment
- PPO: An efficient, stable, and easy-to-implement algorithm
- On-Policy Algorithm works policy to policy
- Clipped surrogate objective function used to strike a balance between exploration and exploitation

Why not SAC?

- An off-policy algorithm learns from a variety of experiences
- Efficient exploration due to entropy maximization
- More difficult to implement than PPO
- Less computationally efficient than PPO due to replay buffer and dual networks

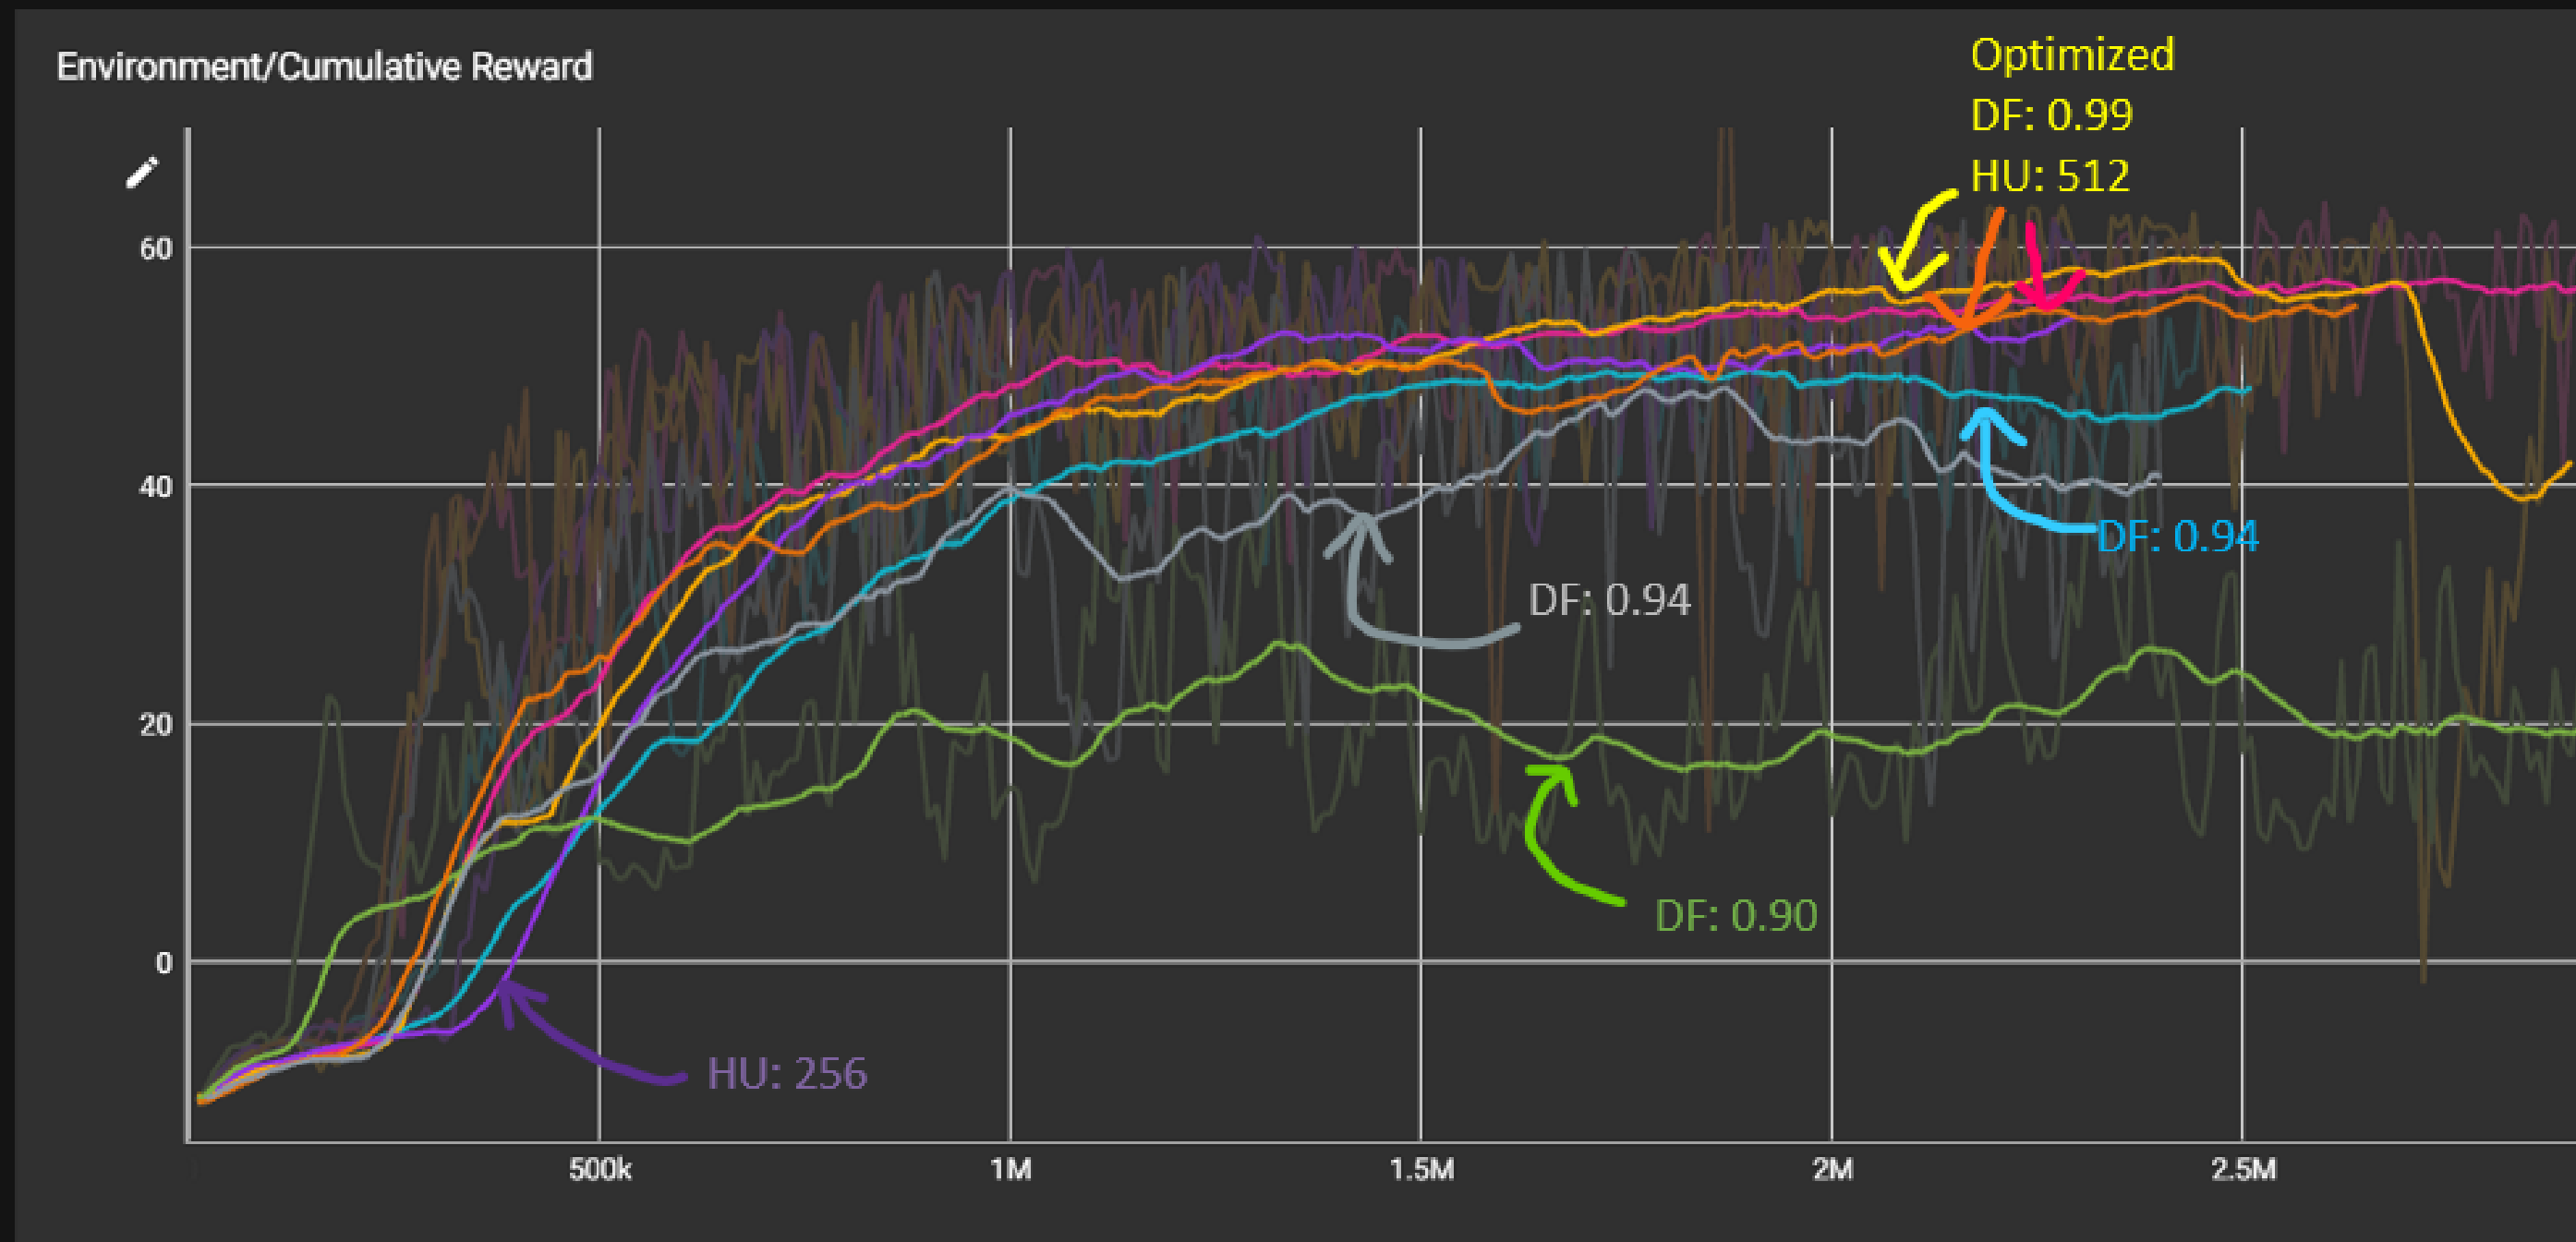
Project Architecture (RL MDP)



Tuning


- Hyperparameters – Discount factor, Learning rate, Epsilon
- Reward scaling – Adjusting the magnitude of rewards to improve training stability
- Identifying bottlenecks – Addressing areas where agents struggle or fail
- Monitoring training progress – Tracking cumulative rewards, loss values, and episode lengths

Results



Conclusion

- Hyperparameters – Discount factor, Learning rate, Epsilon
- Reward scaling – Adjusting the magnitude of rewards to improve training stability
- Identifying bottlenecks – Addressing areas where agents struggle or fail
- Monitoring training progress – Tracking cumulative rewards, loss values, and episode lengths



**Do you have
any questions?**